

Threading and Multiprocessing

Tom Dalton

GitHub: <http://bit.ly/1Dqlw5s>

Parallelising Problems

Modern commodity computer hardware good at parallel computation

- PC Processors
- Graphic Cards

Why do stuff in Parallel?

- Some problems naturally structured this way
 - E.g. Calculate first 100 square numbers
- Can reduce the *real* time taken to get a result

Who not do stuff in Parallel?

- Some problems do not naturally parallelise
 - Calculating fibonacci sequence
- It's Complicated!
 - Synchronisation

Threading

Lightweight parallelism supported by the OS

- share process memory
- separate processor

Threading

File Edit View Search Terminal Help

```
top - 13:11:20 up 3:50, 4 users, load average: 1.80, 0.59, 0.28
Tasks: 223 total, 1 running, 221 sleeping, 0 stopped, 1 zombie
%Cpu(s): 78.8 us, 2.0 sy, 0.0 ni, 19.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 8082836 total, 4126952 used, 3955884 free, 90824 buffers
KiB Swap: 8294396 total, 0 used, 8294396 free. 2024336 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	%CPU	MEM	TIME+	COMMAND
20185	tom	20	0	1081188	28004	11740	314.4	0.3	0:06.93	python
1355	root	20	0	646592	229236	207764	4.0	2.8	5:09.39	Xorg
2470	tom	20	0	491580	32476	13624	4.0	0.4	1:45.64	compiz
2997	tom	20	0	889616	32616	14900	2.0	0.4	0:19.85	gnome-terminal

Multiprocessing

Separate processes can run together

- Often supported by OS
 - Shared memory, pipes, semaphores etc

In the context of this talk

- Working together for a common goal

Multiprocessing

File Edit View Search Terminal Help

```
top - 13:14:58 up 3:53, 4 users, load average: 0.69, 0.61, 0.36
Tasks: 230 total, 6 running, 223 sleeping, 0 stopped, 1 zombie
%Cpu(s): 99.0 us, 0.7 sy, 0.0 ni, 0.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 8082836 total, 4351416 used, 3731420 free, 92744 buffers
KiB Swap: 8294396 total, 0 used, 8294396 free. 2107944 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
21373	tom	20	0	437876	17060	1104	R	94.4	0.2	0:04.82	python
21374	tom	20	0	437876	17064	1108	R	94.4	0.2	0:04.46	python
21372	tom	20	0	437876	17056	1100	R	93.4	0.2	0:04.45	python
21375	tom	20	0	437876	17052	1096	R	93.4	0.2	0:04.48	python
21368	tom	20	0	660252	31032	12088	S	14.9	0.4	0:04.25	python
2470	tom	20	0	491784	32620	13624	S	2.0	0.4	1:50.37	compiz
1355	root	20	0	660900	244432	222620	S	2.0	3.0	5:20.23	Xorg

Example Problems

1. File Search
2. Fractal Plotting

Problem 1: File Search

Find regexp matches in a file

Problem 1: File Search

Quantitative comparisons

- Test file 25M lines (~1.1GB)
- Lines of 10-80 chars
- Each char in [a-z] or space

Problem 1: File Search

Find all regexp matches in a file

```
~/python_edinburgh/parallelism_talk master > grep -n "abcde[a-m]" big_sample_file
211365:kxalvvufy ymqh aaa rplyjqdc bragfirzhwsrvibbxzabcdeirgsjxvn mksecvaswjt xeim
1112241:dlympvzxkcleabcdefobjglpcg igxg qvsx kuwxrvqnpzmzmekoifhbzwingd hkkklk
2640629:bqsvvid zxv rx fsowabcdehy ag vydpbzbfb
3523191:lho rusezzx rlhbg qxwjlxbxm az ezmmqoa edto abcdeetqf
4060083:sitdaqixuczj sptnztpyaozmbtu kazdeg gaabcdehmr araanosh e cfbnxqxmzcqz
4299996:mmkmhyzfecuyboicaphv eogljualf pdlhvabcdedawmszfzbfexkiankz chmrrr kwcsrtchd
4350217:glrxjwufmvdxeaanzegvknnddlcfmfihlqnsdmphczsababcdekciwtudv opdwsqcf
6461019:x bkkicd pxm fwegjqylpsgabcdeibg ybmkf
7039561:jtgshvtvznm mngxueimhyiwtntdsgwpxabcdehwnk
9105581:nh abcdeh gkjmigoj z
10711291:ny wy fqpobgszflbdjhukmabcdeiiqlundayx elkc ii wq
10879350:dxeapeaqvlfr exhbh d omufi nfoilctvpu mhzidljf nihmtf wmem fo abcdehf
12363608:oa xprgf fjkabcdehqvlmid
13986724:ovfgn hricppdqnaabcdeadaq rfr vwzbfaorvt yoqtily
15527376:pyddudntqdrjeqlfccabcdee fllwdfntkrmmakn bztlwe qthlio
18075538:nfgcjabcdeaeecdykmgk
18802776:bx tm iwozcstfahgskbymmjfeut ws sklpouq qmo nabcdebljd pyhoayjyrh
20627195:d kygjmsrq xecr ibnsabcdefwbehbbsoewglrtvc nfcxrma stah ib evmvt
22332650:zamwf abcdehuxgjtgrmesblhnunrobad anxvjasddmpnqt kxxdl
~/python_edinburgh/parallelism_talk master >
```

Problem 1: File Search

Trivial solution

```
for each line in the file:  
    if line matches regexp:  
        output line number and line
```

Problem 1: File Search

```
import re

class FileSearcher(object):

    def __init__(self, filename, regexp_str):
        self.filename = filename
        self.regexp = re.compile(regexp_str)

    def find_matches(self):
        results = {}
        with open(self.filename, "rb") as f:
            for i, line in enumerate(f):
                matchobj = self.regexp.match(line)
                if matchobj:
                    results[i] = line
        return results
```

Problem 1: File Search

Results (Elapsed real time):

Algorithm	Regex matches no lines	Regex matches all lines
Trivial	~10s	~14s

Problem 1: File Search

Parallelised solutions (1)

Read file into chunks

 regex-match chunks in parallel

Problem 1: File Search

Results (Elapsed real time):

Algorithm	Regex matches no lines	Regex matches all lines
Trivial	~10s	~14s
Threaded Regex	~21s	~27s
Multiprocess Regex	~30s	~60s

Problem 1: File Search

Why are these results so bad?

Threading:

- GIL

Multiprocessing:

- multiprocessing.Queue -> Pickle

Problem 1: File Search

What can we do about these problems?

GIL

- Use multiprocessing (!)
- C extensions (more later)

Problem 1: File Search

What can we do about these problems?

Pickle

- Avoid pickling large data sets
- Use shared memory or other IPC

Problem 1: File Search

Parallelised solutions (2)

Every thread/process reads the file

Calculates its own chunk and processes it

Problem 1: File Search

Results (Elapsed real time):

Algorithm	Regex matches no lines	Regex matches all lines
Trivial	~10s	~14s
Threaded Regex	~21s	~27s
Multiprocess Regex	~30s	~60s
Threaded Chunking	~26s	~37s

Problem 1: File Search

Results (Elapsed real time):

Algorithm	Regex matches no lines	Regex matches all lines
Trivial	~10s	~14s
Threaded Regex	~21s	~27s
Multiprocess Regex	~30s	~60s
Threaded Chunking	~26s	~37s
Multiprocess Chunking	~6s	~37s

Problem 1: File Search

Results (Elapsed real time):

Algorithm	Regex matches no lines	Regex matches all lines
Trivial	~10s	~14s
Threaded Regex	~21s	~27s
Multiprocess Regex	~30s	~60s
Threaded Chunking	~26s	~37s
Multiprocess Chunking	~6s	~37s
Grep	~2s	~3s

Problem 1: File Search

Conclusion

For simple problems, it can be hard to beat the overheads

Problem 1: File Search



Problem 2: Plotting Fractals

Mandelbrot fractal 101

- Each point represents a complex number
- Points may diverge or not
- Because maths!
- Pixel colour relates to divergence speed

TLDR Each pixel requires doing a hard sum

Problem 2: Plotting Fractals

[illegible]

<http://commons.wikimedia.org/wiki/File:Mandel.png>

Problem 2: Plotting Fractals

The baseline/control

- Single process/thread
- Calculates each pixel value
- Pixel calc implemented in pure python

Problem 2: Plotting Fractals

Results (Elapsed real time):

Algorithm	Fractal Calc	Iterations	Render Time
Simple	Python	256	~26s

Problem 2: Plotting Fractals

Results (Elapsed real time):

Algorithm	Fractal Calc	Iterations	Render Time
Simple	Python	256	~26s
Threaded	Python	256	~41s
Multiprocess	Python	256	~15s

THREADING



Y U SO SLOW

Problem 2: Plotting Fractals

Threaded slowest!

- GIL

Problem 2: Plotting Fractals

What can we do about these problems?

GIL

- Use multiprocessing
- C extensions

Problem 2: Plotting Fractals

```
int _calc_point(double x0, double y0, int max_iterations) {  
    int i;  
    double x_temp;  
    double x = 0.0;  
    double y = 0.0;  
    for (i=0; i < max_iterations; i++) {  
        if (x * x + y * y >= 4.0) {  
            return i;  
        }  
        x_temp = x * x - y * y + x0;  
        y = 2 * x * y + y0;  
        x = x_temp;  
    }  
    return -1;  
}
```

Problem 2: Plotting Fractals

```
static PyObject * calc_point(PyObject *self, PyObject *args) {  
    double x0, y0;  
    int i, max_iterations;  
    if (!PyArg_ParseTuple(args, "ddi", &x0, &y0, &max_iterations)) {  
        return NULL;  
    }  
    Py_BEGIN_ALLOW_THREADS;  
    i = _calc_point(x0, y0, max_iterations);  
    Py_END_ALLOW_THREADS;  
    if (i == -1) {  
        Py_RETURN_NONE;  
    } else {  
        return Py_BuildValue("i", i);  
    }  
};
```

Problem 2: Plotting Fractals

```
static PyMethodDef c_mandelbrot_methods[] = {  
    {"calc_point",  calc_point, METH_VARARGS, "Calculate a point on a mandelbrot fractal."},  
    {NULL, NULL, 0, NULL}  
};  
  
PyMODINIT_FUNC initlemandelbrot(void) {  
    (void) Py_InitModule("cmandelbrot", c_mandelbrot_methods);  
};
```

Problem 2: Plotting Fractals


Results (Elapsed real time):

Algorithm	Fractal Calc	Iterations	Render Time
Simple	C (optimised)	20000	~26s
Threaded	C (optimised)	20000	~17s
Multiprocess	C (optimised)	20000	~12s

Problem 2: Plotting Fractals

Multiprocess CPU Graph

Render Start



Problem 2: Plotting Fractals

Final Optimisations

- Line-centric
- Random render order

Problem 2: Plotting Fractals

Results (Elapsed real time):

Algorithm	Fractal Calc	Iterations	Render Time
Multiprocess	C (optimised)	20000	~12s
Threaded (line)	C (optimised)	20000	~10s
Multiprocess (line)	C (optimised)	20000	~10s
Threaded (line, random)	C (optimised)	20000	~9s
Multiprocess (line, random)	C (optimised)	20000	~8s

Problem 2: Plotting Fractals

Results (Elapsed real time):

Algorithm	Fractal Calc	Iterations	Render Time
Multiprocess	C (optimised)	20000	~12s
Multiprocess (line)	C (optimised)	20000	~10s
Multiprocess (line, random)	C (optimised)	20000	~8s
Multiprocess (line, random)	Python	20000	~819s

Conclusions

Threading

- Simple
- Lots of Caveats

Multiprocessing

- More complex/less performant IPC
- Overall faster for both examples

WAT

<https://www.destroyallsoftware.com/talks/wat>

- Gary Bernhardt
- CodeMash 2012

WAT

```
#!/usr/bin/env python
from multiprocessing import Process, Queue

def thing(queue, x):
    result = x * 2
    queue.put(result)

q = Queue()
p = Process(target=thing, args=(q, 10, ))

p.start()

p.join()
print q.get()
```

WAT

```
> ./wat.py
```

20

WAT

```
#!/usr/bin/env python
from multiprocessing import Pool, Queue

def thing(queue, x):
    result = x * 2
    queue.put(result)

q = Queue()
pool = Pool(1)

print pool.apply(thing, (q, 10, ))
```

WAT

```
› ./wat2.py
```

```
Traceback (most recent call last):
```

```
  File "./wat2.py", line 11, in <module>
```

```
    print pool.apply(thing, (q, 10, ))
```

```
  File "/usr/lib/python2.7/multiprocessing/pool.py", line 244, in apply
```

```
    return self.apply_async(func, args, kwds).get()
```

```
  File "/usr/lib/python2.7/multiprocessing/pool.py", line 558, in get
```

```
    raise self._value
```

```
RuntimeError: Queue objects should only be shared between processes through inheritance
```


WAT.



Questions?

Thanks!

Slides and Code:

<http://bit.ly/1Dqlw5s>

<https://github.com/tom-dalton-fanduel/python-parallelism-talk>

Parallelism vs Concurrency:

<http://yosefk.com/blog/parallelism-and-concurrency-need-different-tools.html>